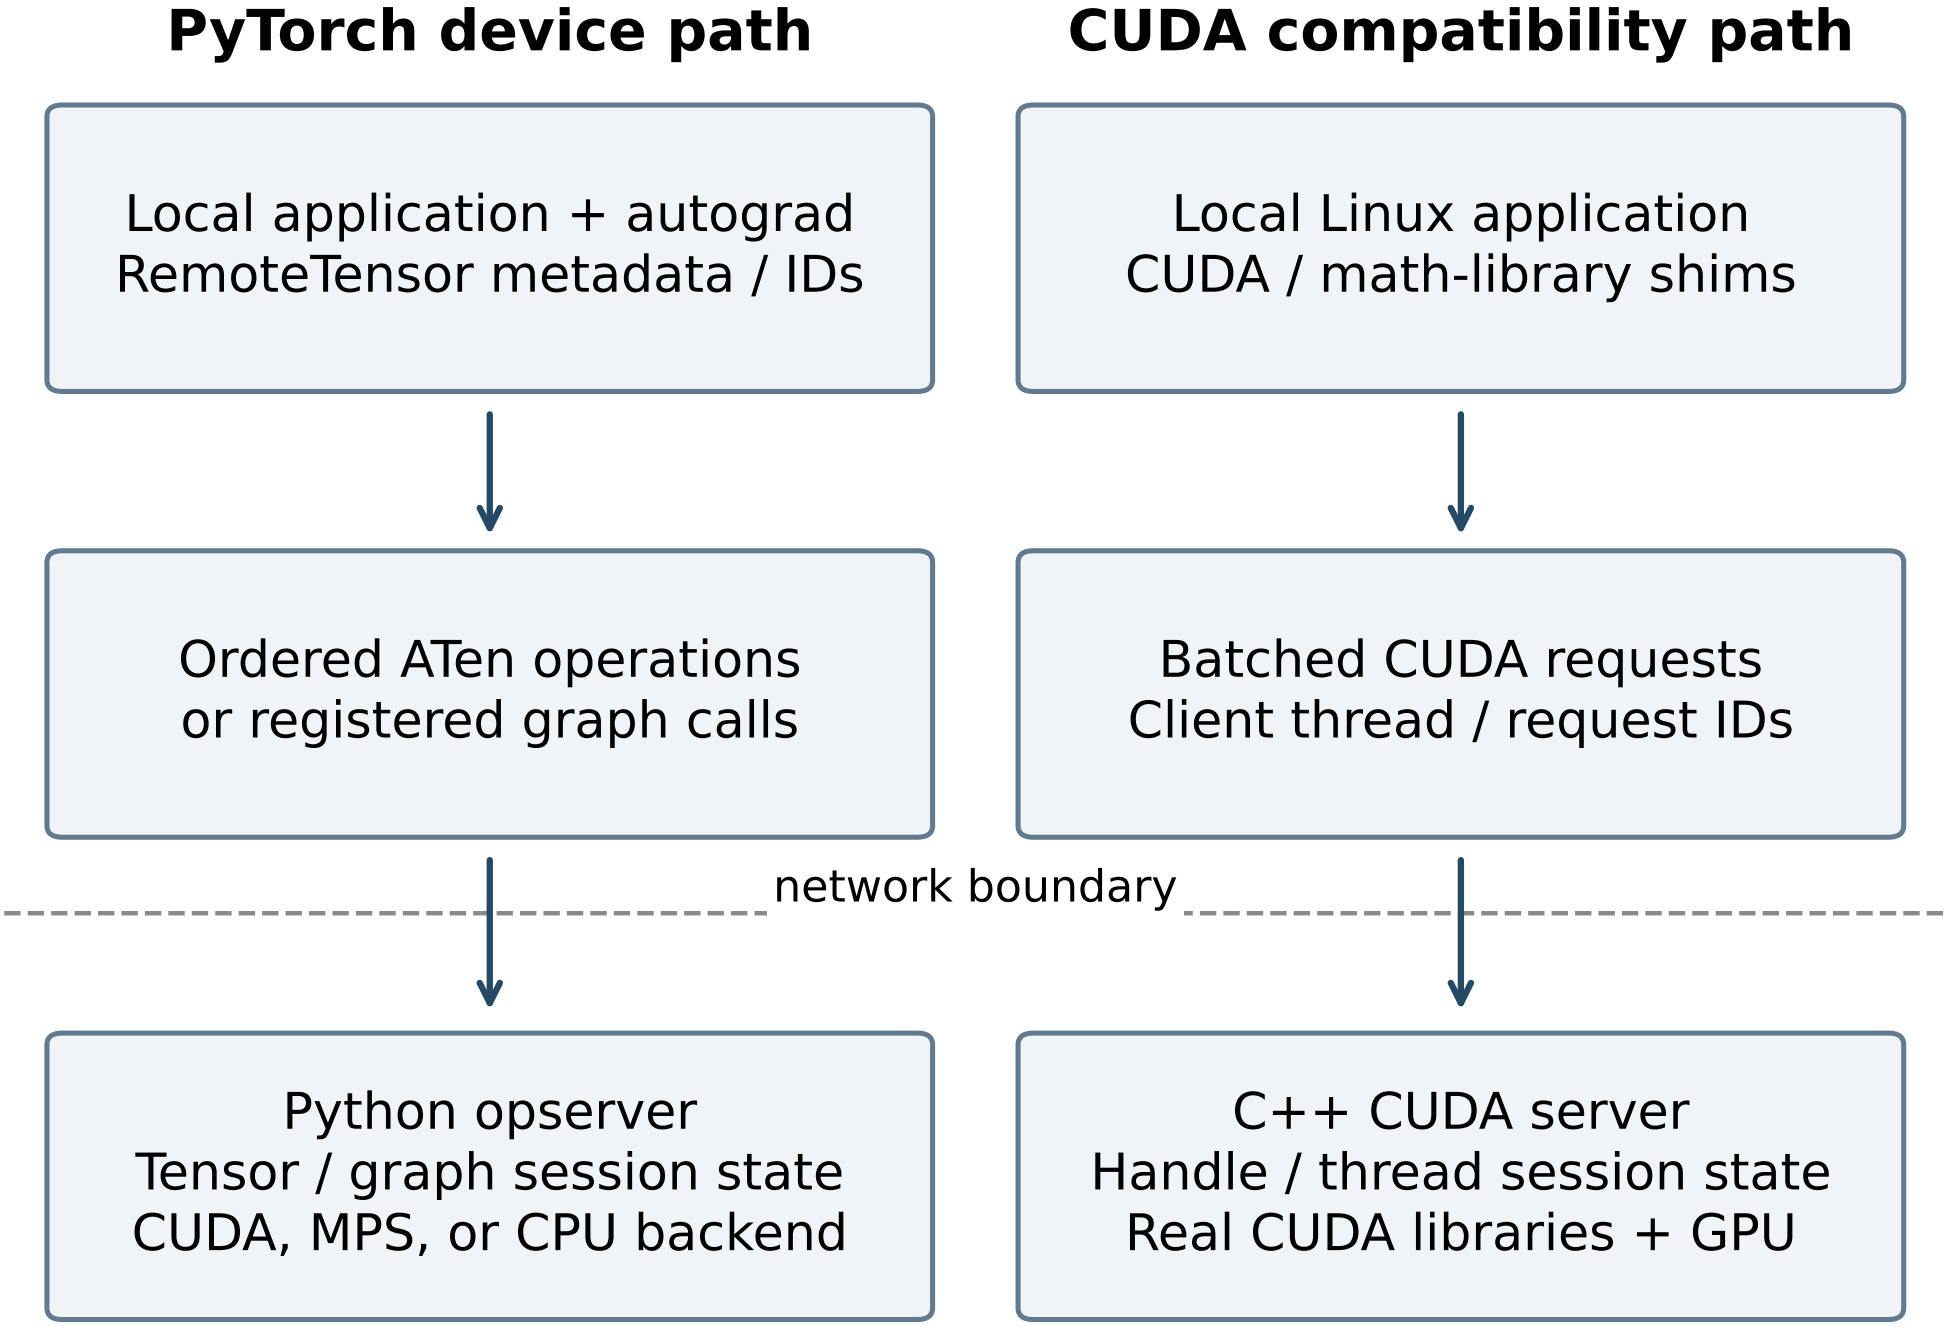
PyTorch device path
CUDA compatibility path
Local application + autograd
RemoteTensor metadata / IDs
Local Linux application
CUDA / math-library shims
Ordered ATen operations
or registered graph calls
Batched CUDA requests
Client thread / request IDs
network boundary
Python opserver
Tensor / graph session state
CUDA, MPS, or CPU backend
C++ CUDA server
Handle / thread session state
Real CUDA libraries + GPU